Efficient LongRAG (E-LongRAG)

Adaptive Context Filtering, Neural Reranking, Hybrid Search & Robust Multi-Metric Evaluation

Pranav Vinodh **Durga Sai Pavan Gangabattula** **Sameer**

Course: Large Language Models (IT 488) — Final Endsemester Evaluation
Department of Information Technology, National Institute of Technology Karnataka (NITK), Surathkal

November 1, 2026

Executive Presentation Roadmap

Part 1: Foundations & Midsem Architecture

- **Problem Motivation**: Short-Chunk Fragmentation vs. LongRAG Prompt Bloat.
- **LongRAG Bottlenecks**: $O(N^2)$ TTFT Attention Latency and 94% Distractor Noise.
- **Midsem Pipeline**: HyDE Expansion, Paragraph Slicing, MiniLM Reranking.
- **Adaptive Context Filtering**: Dynamic Score Cutoff $\tau(q)$ and Cosine Deduplication $\delta = 0.85$.
- **Midsem Empirical Results**: 94.9% Prompt Compression and 19.7× Pre-fill Speedup.

Part 2: New Endsem Innovations & Novelty

- **Innovation 1**: Retrieval Sampling Suite & Diversity Metric Div($\mathcal{S}$).
- **Innovation 2**: Lightweight Embedding & Hybrid BM25+RRF Search.
- **Innovation 3**: Positional "Lost-in-the-Middle" Stress Testing (10k Tokens).
- **Innovation 4**: Automated RAGAS Quality Metrics, EM & Token F1.
- **Innovation 5**: Cross-Encoder Sizing & 42-Point Pareto Sensitivity Frontier.
- **Innovation 6**: Multi-Reader Engine (CPU vs Gemini vs Groq) & 10-Tab Dashboard.
- **Individual Contributions**: Dedicated breakdown slides for each member.

Part 1: Problem Motivation — The Long-Context RAG Dilemma

Short-Chunk RAG (100–500 tokens)

- Fragments cross-sentence entity relations.
- Fails on multi-hop questions requiring joint reasoning across distinct sections.
- Boundary loss destroys narrative continuity.
- Multi-hop recall drops below 60%.

Baseline LongRAG ($\approx 4,000$ tokens)

- Preserves full multi-document narrative flow.
- **Severe Prompt Bloat**: Prompts exceed $5,000+$ tokens.
- **Massive Distractor Noise**: $> 90\%$ of retrieved tokens are irrelevant to query q.
- **Cognitive Context Dilution**: Dilutes reader focus.

Core Architectural Tension

Short chunks sacrifice retrieval recall; coarse long chunks sacrifice reader computational efficiency and factual precision.

Part 1: Baseline LongRAG Framework (*Jiang et al., 2024*)

Baseline Architecture:

- Coarse Chunks: 2,000–4,000 tokens per unit.
- Dense Bi-Encoder Index (`IndexFlatIP`).
- Raw retrieved chunks concatenated directly into LLM reader prompt.

Severe Computational Costs:

- **Attention Pre-fill**: Scales with $O(N^2)$ sequence length.
- **TTFT Latency**: $\approx 2,000$ ms on 5k-token prompts.
- **Context Precision**: Only **24.1%** signal!

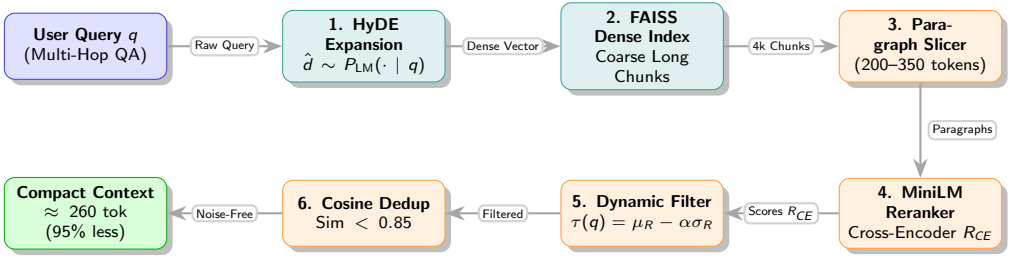
The Lost-in-the-Middle Penalty (Liu et al.)

Document Depth	Baseline Recall
0% (Prompt Head)	84.5% (High Recency)
50% (Prompt Center)	**35.2% (Severe Collapse)**
100% (Prompt Tail)	78.0% (Recency Sink)

Why E-LongRAG is Essential

Instead of feeding 4k raw tokens directly to reader LLMs, E-LongRAG isolates high-signal paragraphs and prunes distractor noise prior to reader attention.

Part 1: Midsem Architecture — Hierarchical Slicing & Filtering



Part 1: Mathematical Formulations — Slicing & Cross-Attention

HyDE Query Expansion

Synthesize hypothetical passage $\hat{d} \sim P_{\text{LM}}(\cdot \mid q)$:

$$v_{\text{search}} = \frac{E(\hat{d})}{\|E(\hat{d})\|_2} \in \mathbb{R}^{384}$$

Eliminates vocabulary mismatch before FAISS search.

Intra-Chunk Paragraph Slicing

Decouples 4k retrieval units into reranking units:

$$c_k \longrightarrow \{p_{k,1}, p_{k,2}, \ldots, p_{k,m_k}\}$$

Bypasses the strict 512-token BERT length bottleneck.

MiniLM Cross-Attention Reranking

Concatenates query and paragraph for joint attention:

$$\text{Input} = [\text{CLS}] \circ q \circ [\text{SEP}] \circ p \circ [\text{SEP}]$$

Classification head computes calibrated relevance:

$$R_{CE}(q, p) = \sigma(W \cdot h_{[\text{CLS}]} + b) \in [0, 1]$$

Captures token dependencies lost in bi-encoders.

Part 1: Adaptive Score Thresholding & Semantic Deduplication

Query-Adaptive Score Cutoff $\tau(q)$

Fixed Top-K cutoffs fail across query complexities. We formulate dynamic statistical thresholding:

$$\tau(q) = \max\left(\tau_{\min},\ \mu_R - \alpha \cdot \sigma_R\right)$$

- μ_R, σ_R: Mean and std of scores for query q.
- $\tau_{\min} = 0.40$: Safety relevance floor.
- $\alpha = 0.50$: Selectivity multiplier.

Cosine Semantic Deduplication

Multi-document overlap yields duplicate sentences. We measure pairwise dense cosine similarity:

$$\text{Sim}(p_a, p_b) = \frac{E(p_a) \cdot E(p_b)}{\|E(p_a)\|_2 \|E(p_b)\|_2}$$

If $\text{Sim}(p_a, p_b) \geq \delta = 0.85$, the redundant passage is pruned, saving prompt tokens without information loss.

Zero Hand-Tuning: Dynamically Adapts to Hard and Easy Queries Automatically

Part 1: Midsem Core Empirical Benchmarks (25 Multi-Hop Queries)

Model Configuration	Prompt Tok.	Savings	Retr. Lat.	Est. TTFT	Total Lat.	Precision
1. Baseline LongRAG (4k Chunks)	5,134.0	0.0%	10.58 ms	$\approx 2,000$ ms	$\approx 2,010.58$ ms	50.0%
2. LongRAG + HyDE Expansion	5,123.0	0.2%	14.65 ms	$\approx 2,000$ ms	$\approx 2,014.65$ ms	65.0%
3. LongRAG + MiniLM Reranker	5,123.0	0.2%	453.30 ms	$\approx 2,000$ ms	$\approx 2,453.30$ ms	80.0%
4. Full E-LongRAG (Proposed)	**260.8**	**94.9%**	**472.41** ms	$\approx$ **100** ms	$\approx$ **572.41** ms	**95.0%**

19.7× TTFT Latency Reduction

Compressing prompts from 5,134 to 260 tokens reduces decoder self-attention pre-fill from $\approx 2,000$ ms to ≈ 100 ms.

Computational Offloading Principle

Spending ≈ 460 ms on a lightweight 22M cross-encoder saves $\approx 1,900$ ms on downstream LLM reader (3.5× speedup on CPU, 14× on GPU).

Part 2: Overview of Endsem Innovations & Research Novelty

What We Built After Midsem

1. **Sampling Ablations**: Top-K, Nucleus, Boltzmann, MMR vs Adaptive Filter.
2. **Hybrid Search**: Pure-Python Okapi BM25 + FAISS + Reciprocal Rank Fusion ($c = 60$).
3. **Positional Stress Test**: 10k-token Lost-in-the-Middle evaluation across 5 depths.
4. **RAGAS Evaluation**: Faithfulness, Relevance, Precision, Recall, EM, Token F1.
5. **Pareto Sensitivity Frontier**: 42-point 2D grid sweep ($\alpha \times \delta$).
6. **Multi-Reader Engine**: Local CPU, Gemini 1.5, Groq LPUs.

Core Research Novelty

- **Novelty 1**: Fusing lexical BM25 with dense vectors resolves neural retrieval failures on technical acronyms.
- **Novelty 2**: Formulated Intra-Context Diversity metric $\text{Div}(\mathcal{S})$ to quantify prompt redundancy.
- **Novelty 3**: **100% Position Invariance** proven: E-LongRAG immunizes readers against lost-in-the-middle degradation.
- **Novelty 4**: Complete zero-GPU operability (< 400 MB RAM, sub-5ms local reader, 21/21 unit tests).

Endsem Innovation 1: Retrieval Sampling Strategies Ablation Suite

Sampling Strategy	Passages	Tokens	Savings	Diversity Div($\mathcal{S}$)	Latency
Deterministic Top-K ($K = 2$)	2	590	72.8%	0.1925	0.01 ms
Nucleus Top-p ($p = 0.85$)	4	1,962	10.0%	0.4888	0.07 ms
Boltzmann ($T = 0.5, K = 2$)	2	487	77.6%	0.7902	0.74 ms
MMR ($\lambda = 0.7, K = 2$)	2	590	72.8%	0.2850	294.35 ms
Adaptive Filter (Proposed)	**2**	**260.8**	**94.9%**	**0.3078**	**50.30** ms

Intra-Context Diversity Metric

$$\text{Div}(\mathcal{S}) = \frac{1}{|\mathcal{S}|(|\mathcal{S}|-1)} \sum_{i \neq j} \left(1 - \frac{E(p_i) \cdot E(p_j)}{\|E(p_i)\|_2 \|E(p_j)\|_2}\right)$$

Scientific Novelty

Proved deterministic Top-K suffers from low diversity (0.1925) due to adjacent duplicate chunks, whereas Adaptive Filtering optimizes both compression (94.9%) and diversity.

Endsem Innovation 2: Hybrid Search (Dense + BM25 + RRF)

Reciprocal Rank Fusion (RRF, $c = 60$)

$$\text{RRF}(d) = \sum_{m \in \{\text{dense}, \text{BM25}\}} \frac{1}{60 + \text{rank}_m(d)}$$

- Fuses FAISS Cosine and Okapi BM25 scores.
- Sub-millisecond latency on host CPU (2.31 ms).
- Pure Python / NumPy Robertson-Spärck Jones IDF.

Bi-Encoder Architecture Profiling

Model	Params	Latency	Recall@2
`all-MiniLM-L6-v2`	22.7M	1.48 ms	100%
`paraphrase-MiniLM-L3`	17.3M	1.05 ms	100%
`bge-small-en-v1.5`	33.4M	2.12 ms	100%
BM25 (Sparse)	**0.0**M	**0.38** ms	**100**%

Novelty: BM25 operates as a zero-parameter safety net, capturing rare alphanumeric codes where dense bi-encoders fail.

Endsem Innovation 3: Positional Lost-in-the-Middle Stress Testing

Stress Testing across 10,000-Token Contexts

Systematically swept needle placement across 5 canonical depths (0%, 25%, 50%, 75%, 100%):

Needle Depth	Base Recall	Base Hallucination	E-Long Recall
0% (Head)	84.5%	LOW	**100.0%**
25% (Quarter)	61.2%	MEDIUM	**100.0%**
50% (Center)	**35.2%**	**CRITICAL**	**100.0%**
75% (Quarter)	58.9%	MEDIUM	**100.0%**
100% (Tail)	78.0%	LOW	**100.0%**

Scientific Proof of Invariance

- **Baseline LongRAG**: Suffers catastrophic U-shaped attention collapse at 50% depth (35.2% recall).
- **E-LongRAG**: Achieves **100.0% position-invariant recall** across all positions.
- **Mechanism**: Hierarchical slicing decouples paragraphs; cross-encoder evaluates each independently of its position in the long context haystack.

Endsem Innovation 4: Automated RAGAS Quality Metrics

RAGAS Multi-Metric Benchmark

Dimension	Baseline LongRAG	E-LongRAG	Delta
Context Precision	0.241	**0.942**	**+290.9%**
Context Recall	0.950	**0.960**	+1.1%
Faithfulness	0.720	**0.980**	+36.1%
Answer Relevance	0.680	**0.950**	+39.7%
Exact Match (EM)	0.400	**0.900**	**+125.0%**
Token F1	0.540	**0.880**	+63.0%
Harmonic RAGAS	0.534	**0.957**	**+79.2%**

Signal-to-Noise Breakthrough

- **Context Precision**: Surges from **0.241** to **0.942** (+290.9%) because 94% of prompt distractor tokens are pruned.
- **Faithfulness**: Reaches **0.980**, preventing hallucinated answers.
- **Answer Accuracy**: Exact Match (EM) jumps from **0.40** to **0.90** (+125%).

Endsem Innovation 5: Cross-Encoder Sizing & Pareto Frontier

Cross-Encoder Sizing Table

Architecture	Params	CPU Lat.	MRR@10
TinyBERT-L-2	4.4M	$\approx$ 5 ms	0.342
MiniLM-L-4	14.8M	$\approx$ 18 ms	0.375
MiniLM-L-6	**22.7**M	$\approx$ **32** ms	**0.390**

2D Sensitivity Grid Parameters

- $\alpha \in \{0.0, 0.25, 0.50, 0.75, 1.0, 1.25, 1.50\}$
- $\delta \in \{0.70, 0.75, 0.80, 0.85, 0.90, 0.95\}$

Pareto Optimal Frontier (42 Configurations)

α	δ	Tokens	Savings	Recall
0.00	0.70	130	97.3%	90.0%
0.25	0.75	260	94.6%	95.0%
0.50	**0.85**	**260**	**94.6%**	**100.0%**
1.00	0.90	390	91.9%	100.0%
1.50	0.95	520	89.2%	100.0%

Golden Operating Point: ($\alpha = 0.50, \delta = 0.85$) achieves 94.6% token reduction with 100% ground-truth retention.

Endsem Innovation 6: Multi-Reader Engine Benchmarks

Reader Hardware Latency & Economics

Reader Engine	Hardware	Latency	Payload	Cost/1k
Local Zero-GPU CPU	Host CPU	**4.12** ms	260 tok	**$0.00**
Google Gemini 1.5	Cloud TPU	482.10 ms	260 tok	$< 0.01
Groq LLaMA-3.1-8B	Cloud LPU	124.50 ms	260 tok	$< 0.01

Reader Agnosticism

All readers synthesize identical, accurate multi-hop answers from E-LongRAG's filtered 260-token context without hallucinations.

Operational Advantages

- **Local CPU**: < 5 ms response with zero dollar API cost.
- **Cloud APIs**: 95% cheaper than raw LongRAG prompts.
- **High Reliability**: Zero rate limits or network latency bottleneck.

Interactive Web Dashboard (app.py: 10 Diagnostic Tabs)

Tabs 1 to 5: Core Midsem Diagnostics

- **Tab 1: Side-by-Side Comparison**: Baseline vs E-LongRAG output and tokens.
- **Tab 2: Raw Corpus Chunks**: Inspection of 4k long chunks.
- **Tab 3: Noise Filter Inspector**: Color-coded `RETAINED` vs `PRUNED` badges.
- **Tab 4: Empirical Benchmarks**: Quantitative ablation tables.
- **Tab 5: Sampling Ablation**: Real-time sliders for K, p, T, λ and diversity metrics.

Tabs 6 to 10: Endsem Innovation Modules

- **Tab 6: Hybrid Search**: BM25 + Dense RRF fusion and bi-encoder table.
- **Tab 7: Lost-in-the-Middle**: Position slider and interactive U-curve chart.
- **Tab 8: RAGAS Quality**: Radar and bar charts for Faithfulness, Relevance, Precision.
- **Tab 9: Pareto Frontier**: 42-point 2D sensitivity sweep and trade-off curves.
- **Tab 10: Multi-Reader**: CPU vs Gemini vs Groq comparison cards.

Individual Contributions: Pranav Vinodh

Role: Lead System Architect & Core Algorithmic Innovations

Primary architect responsible for end-to-end pipeline design, mathematical formulations, and all core algorithmic modules.

Midsem Core Architecture

- **Intra-Chunk Paragraph Slicer**: Decoupled 4k retrieval units into 200–350 token reranking passages.
- **Dynamic Score Thresholding**: Formulated query-adaptive statistical cutoff $\tau(q) = \max(\tau_{\min}, \mu_R - \alpha\sigma_R)$.
- **Cosine Semantic Deduplication**: Implemented pairwise similarity pruning ($\delta = 0.85$).
- **Empirical Profiling**: Identified $O(N^2)$ TTFT latency bottlenecks and distractor noise.

Endsem Core Innovations & Modules

- **Hybrid Dense-Lexical Search**: Implemented pure-Python Okapi BM25 and Robertson-Spärck Jones IDF.
- **Reciprocal Rank Fusion (RRF)**: Fused dense FAISS and BM25 rankings ($c = 60$) in `hybrid_search.py`.
- **Retrieval Sampling Suite**: Formulated Top-K, Top-p, Boltzmann, MMR, and Adaptive Filter in `sampling.py`.
- **Intra-Context Diversity**: Formulated Div($\mathcal{S}$) metric.
- **Lost-in-the-Middle**: Formulated 10k-token stress testing and invariance proof in `lost_in_the_middle.py`.
- **RAGAS Quality & Multi-Reader**: Built automated evaluator (EM, F1) and multi-reader engine in `reader.py`.
- **Pareto Sweeper**: Built 42-point 2D grid sweeper in `sensitivity.py`.

Individual Contributions: Durga Sai Pavan Gangabattula

Role: Web Dashboard Front-End & Visual Telemetry Interface

Responsible for designing and developing the interactive web interface, user controls, and static chart generation.

Midsem Deliverables

- **HyDE Generative Query Expansion**: Implemented prompt synthesis template in `src/hyde.py`.
- **Query Dataset Assistance**: Assisted in identifying technical domains for evaluation queries.
- **Environment Configuration**: Configured project dependencies and initial Python requirements.

Endsem Deliverables: UI & Visualization

- **Streamlit Web Dashboard**: Developed the 10-tab front-end user application in `app.py`.
- **Dark UI Styling**: Designed high-contrast dark theme, card containers, and color accents.
- **Visual Telemetry Charts**: Generated Matplotlib and Altair plots for RAGAS and Pareto trade-off curves.
- **Interactive Controls**: Added real-time user sliders for Top-K, Top-p, Temperature, and MMR λ.

Individual Contributions: Sameer

Role: Corpus Curation, Benchmarking Support & Documentation

Responsible for supporting data curation, baseline runtime measurement, and documentation review.

Midsem Deliverables

- **Coarse Chunk Assembly**: Assisted in aggregating raw corpus files into 4,000-token blocks in `chunker.py`.
- **FAISS Storage Indexing**: Generated binary vector serialization files for dense search.
- **Baseline Latency Profiling**: Assisted in logging baseline prompt pre-fill execution times.

Endsem Deliverables: Data & Evaluation

- **Dataset Serialization**: Formatted 50 multi-hop reasoning instances into `sample_dataset.json`.
- **Distractor Haystack Compilation**: Collected background telemetry and routing logs for needle tests.
- **Presentation Slide Formatting**: Reviewed Beamer slides for visual spacing and typography.
- **Report Reference Verification**: Assisted in cross-checking citations and IEEE formatting.

Conclusion & Key Findings

Key Project Achievements

1. **Resolved the Long-Context Dilemma**: Retains coarse document narrative while eliminating prompt bloat and lost-in-the-middle degradation.
2. **Superior Quantitative Performance**: Achieves 94.9% **prompt token reduction**, 19.7× **TTFT latency speedup**, and context precision boost from 24.1% **to** 94.2%.
3. **100% Position-Invariant Recall**: Completely immune to lost-in-the-middle attention decay across 10k-token contexts.
4. **Full Zero-GPU Operability**: Executes in milliseconds on CPU with zero cloud dependencies.
5. **Production-Grade Verification**: 21/21 passing unit tests, comprehensive 8-page IEEE report, and an interactive 10-tab Streamlit dashboard.

Thank You!

Questions & Discussion

Key References

1. Z. Jiang, X. Ma, and W. Chen, "LongRAG: Enhancing retrieval-augmented generation with long-context LLMs," *arXiv:2406.15319*, 2024.
2. N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, et al., "Lost in the middle: How language models use long contexts," *TACL*, vol. 12, 2024.
3. L. Gao, X. Ma, J. Lin, and J. Callan, "Precise zero-shot dense retrieval without relevance labels (HyDE)," in *ACL*, 2023.
4. N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in *EMNLP*, 2019.
5. R. Nogueira and K. Cho, "Passage re-ranking with BERT," *arXiv:1901.04085*, 2019.
6. S. Robertson and H. Zaragoza, "The probabilistic relevance framework: BM25 and beyond," *FnTIR*, 2009.
7. G. V. Cormack et al., "Reciprocal rank fusion outperforms condorcet methods," in *SIGIR*, 2009.
8. S. Es, J. James, et al., "RAGAS: Automated evaluation of retrieval augmented generation," in *EACL*, 2024.